Compressing Deep Networks

James Kwok

DeepLearn 2019

香港科技大學
THE HONG KONG
UNIVERSITY OF SCIENCE
AND TECHNOLOGY

Machine Learning is Everywhere

Self-Driving

Machine Translation

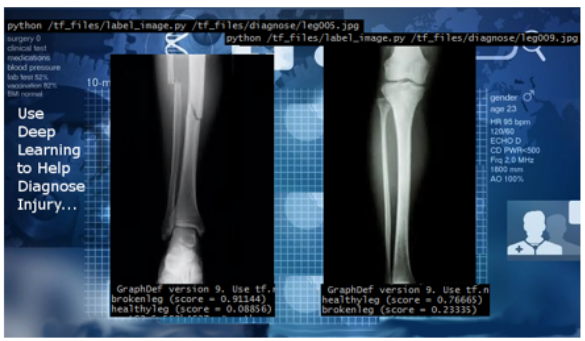

Healthcare

Game Playing

deep learning: excellent performance in a variety of domains

Deep Learning (Neural Networks)

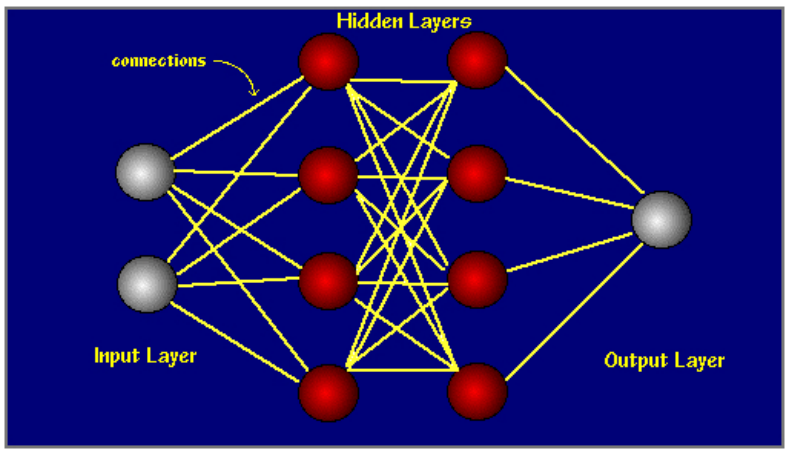

Deeper and Deeper Networks

ImageNet classification

	number of layers	top-5 error (%)
ILSVRC'12 (AlexNet)	8	16.4
ILSVRC'13	8	11.7
ILSVRC'14 (VGG)	19	7.3
ILSVRC'14 (GoogleNet)	22	6.7
ILSVRC'15 (ResNet)	152	3.57

Deep Learning + Big Data + Big Compute

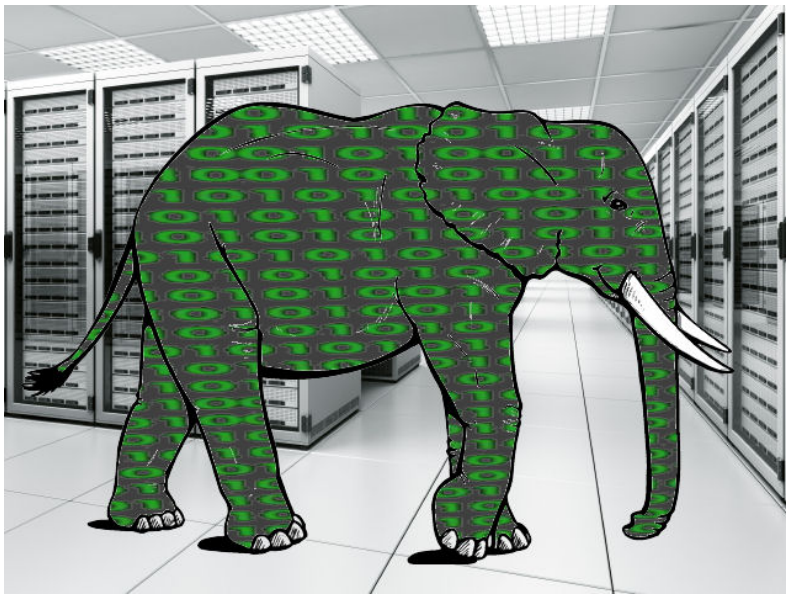

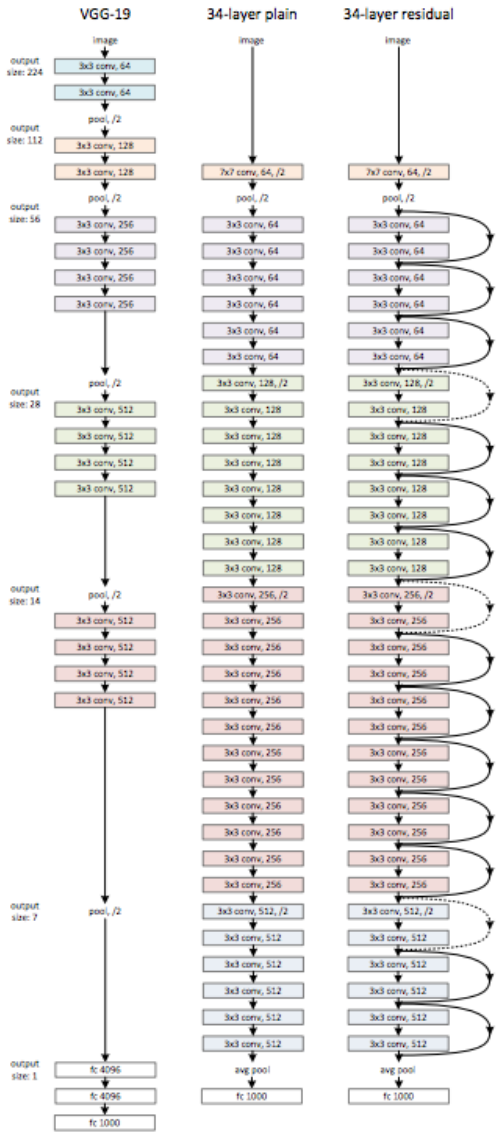

Overfitting

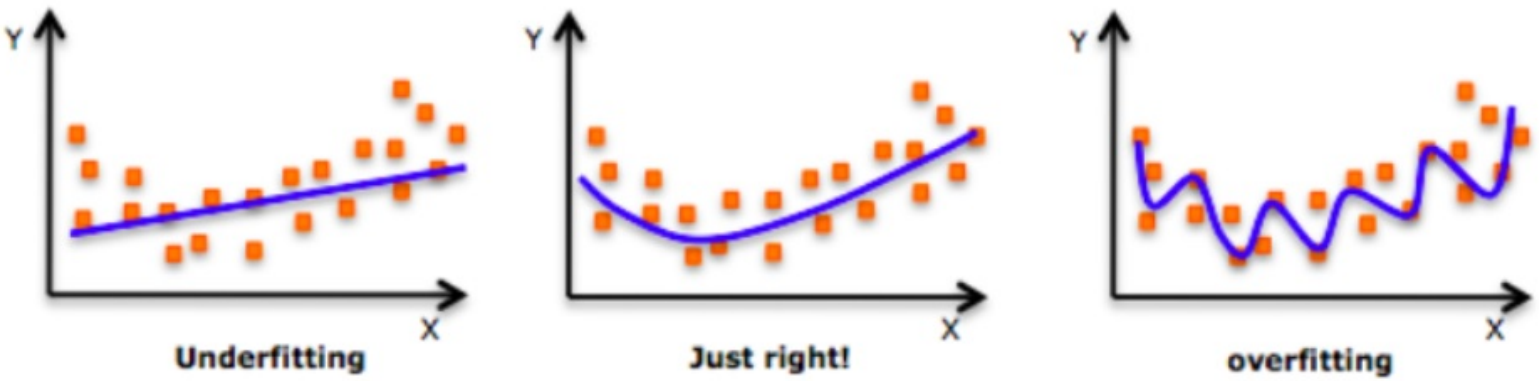

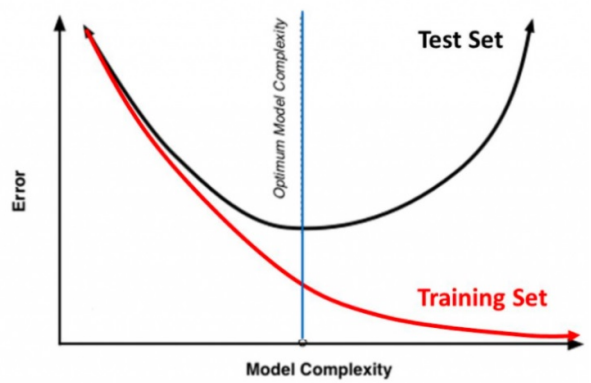

Quest for a Small Model

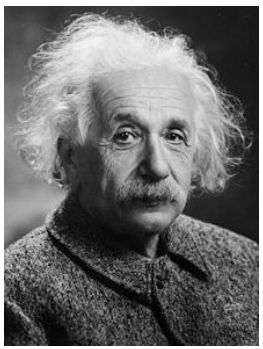

Everything should be made as simple as possible, but no simpler

Occam's razor:
The simplest solution tends to be the right one

Advantages of a small machine learning model

- better generalization
- smaller memory footprint
- faster prediction
- less expensive to collect features

Deep Learning: From Development to Deployment

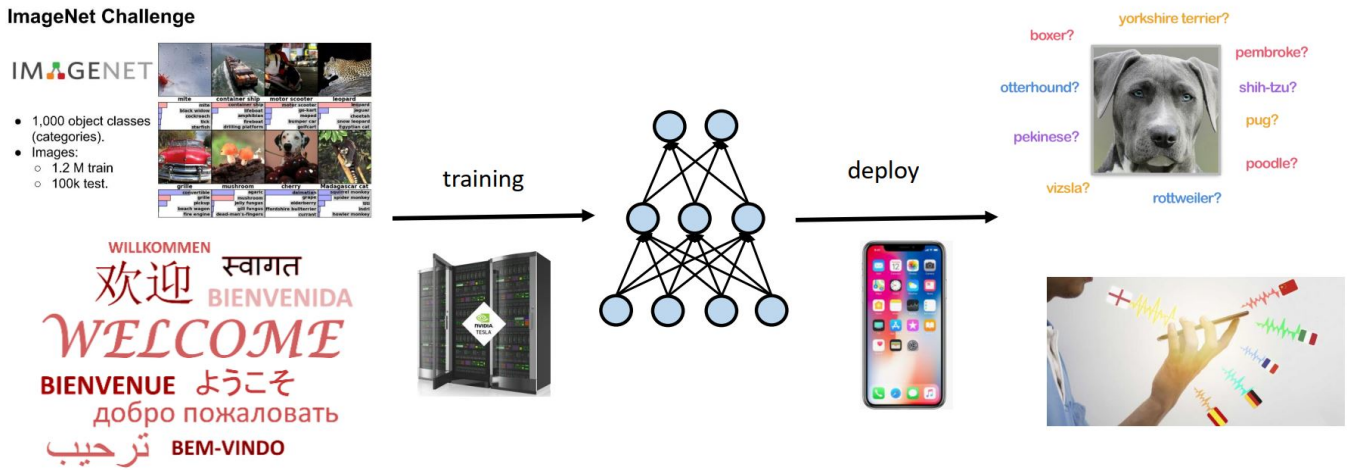

Example (AlexNet, VGG-16, Resnet)

- hundred of megabytes to store
- billions of high-precision operations on classification
- more operations $\rightarrow$ more energy

From Development to Deployment

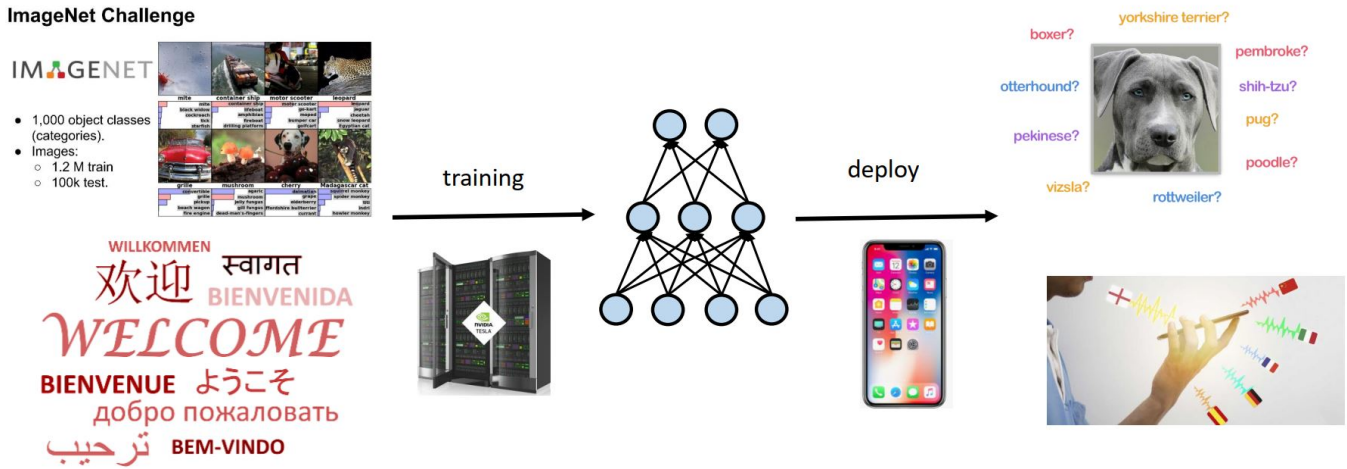

Problem

computation and memory intensive on small computing devices

- cell phones, self-driving cars, internet of things (IoT) devices

Good News!

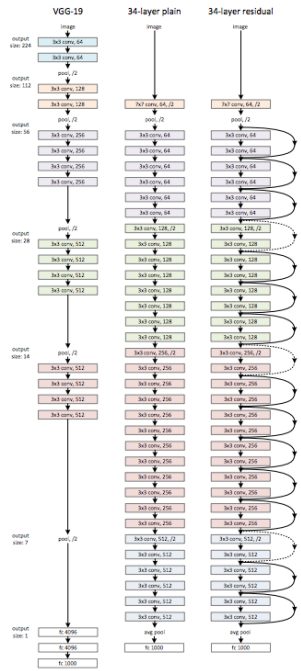

capacity of deep network is usually
larger than needed

can be compressed without accuracy degradation

Compressing Deep Networks

- network sparsification

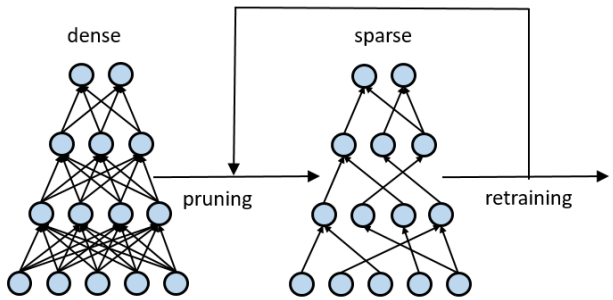

- quantization

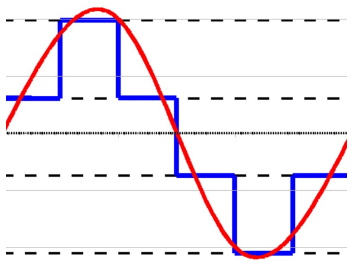

- low-rank approximation

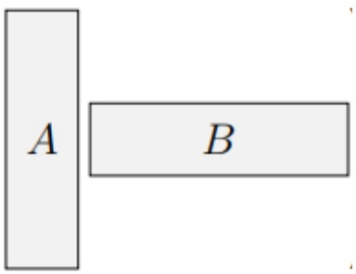

Compressing Deep Networks...

- distillation

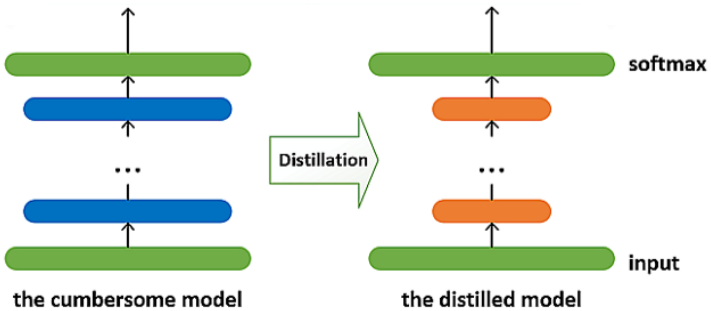

- more compact model

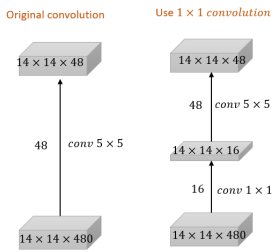

- neural architecture search