

Network Sparsification

reduce redundancy in the **number of weights**

Feature Selection

use only a small subset of the features

Advantages

- better generalization
- smaller memory footprint
- faster prediction
- less expensive to collect features

Challenging

Example

- w : 1000-dimensional; $B = 5$
- number of possible choices: $C_{1000}^5 = 8.25 \times 10^{12}$ (expensive)

heuristic methods

- forward selection
 - the best single-feature is picked first
 - add the next best feature, ...
 - $\{\} \rightarrow \{A_1\} \rightarrow \{A_1, A_4\} \rightarrow \dots$
- backward elimination
 - repeatedly eliminate the worst feature
 - $\{A_1, A_2, A_3, A_4, A_5, A_6\} \rightarrow \{A_1, A_3, A_4, A_5, A_6\} \rightarrow \{A_1, A_3, A_4, A_6\} \rightarrow \dots$
- ...

Regularization

Training samples $\{(x_1, y_1), \dots, (x_n, y_n)\}$

empirical risk minimization

$$\min_w \text{loss}$$

prior knowledge about the model \rightarrow regularizer

regularized risk minimization

$$\min_w \text{loss} + \lambda \underbrace{r(w)}_{\text{regularizer}}$$

What Regularizer?



prior knowledge: many model parameters should be small

Example (ℓ_2 -regularizer)

$$r(w) = \|w\|_2^2 = \sum_{i=1}^d w_i^2$$

square loss + ℓ_2 regularizer \rightarrow ridge regression

Model with Large Number of Parameters

prior knowledge: many parameters are not useful



minimize $\text{loss} + \text{sparsity-inducing regularizer}$


Example (ℓ_0 regularizer)

- $\|w\|_0$: number of nonzero elements in w
- feature selection: $\|w\|_0 \leq B$

ℓ_1 Regularizer

ℓ_1 as a surrogate for ℓ_0

$$r(w) = \|w\|_1 = \sum_{i=1}^d |w_i|$$



Example (Lasso (Tibshirani, 1996))

use square loss: $\min \|y - Xw\|^2$ s.t. $\|w\|_1 \leq t$

Pruning of Deep Networks

what to prune?

- 1 unstructured pruning
- 2 structured pruning


how does the pruned network look like?

Unstructured Pruning of Deep Networks

Prune Weights

(Han et al., 2015)

Learning both Weights and Connections for Efficient Neural Networks



three-step process

- 1 learn the connectivity via normal network training
- 2 prune the connections (dense network \rightarrow sparse network)
- 3 **retrain** the network learn the final weights for the remaining sparse connections

Which Connections to Prune?

based on the

- ① **weight magnitude**
 - assumption: small weights are less important
 - all connections with weights below a **threshold** are removed from the network
 - encourage unimportant weights to smaller values
→ usually also use ℓ_1 or ℓ_2 regularization
- ② **hidden unit activation**
 - small activation value → this feature detector is less important for prediction
- ③ **change in loss** after removing weights
- ④ ...

Experimental Results

Lenet on MNIST, AlexNet/VGG-16 on ImageNet

Network	Top-1 Error	Top-5 Error	Parameters	Compression Rate
LeNet-300-100 Ref	1.64%	-	267K	
LeNet-300-100 Pruned	1.59%	-	22K	12×
LeNet-5 Ref	0.80%	-	431K	
LeNet-5 Pruned	0.77%	-	36K	12×
AlexNet Ref	42.78%	19.73%	61M	
AlexNet Pruned	42.77%	19.67%	6.7M	9×
VGG-16 Ref	31.50%	11.32%	138M	
VGG-16 Pruned	31.34%	10.88%	10.3M	13×

[from (Han et al., 2015)]

- network pruning can save 9x to 13x parameters with no drop in predictive performance

Regularizing Deep Network Weights: ℓ_2

(Collins et al., 2014)

Memory Bounded Deep Convolutional Networks

ℓ_2 regularizer

$$R(\mathbf{w}) = \|\mathbf{w}\|_2^2 \text{ (weight decay)}$$


reduce the weight magnitude, but **not** exactly to zero

Regularizing Network Weights: ℓ_1

ℓ_1 regularizer


$$R(\mathbf{w}) = \|\mathbf{w}\|_1$$

- nonsmooth
- replace gradient by **subgradient**



subgradient

- g is a **subgradient** of f at x iff $f(y) \geq f(x) + g'(y - x)$
- may not be unique
- if f is differentiable \rightarrow gradient




subgradient descent

$$x_{t+1} = x_t - \eta g_t, \text{ for some subgradient } g_t \text{ of } f \text{ at } x_t$$

Subgradient Descent for ℓ_1

Example ($f(x) = |x|$)



$$w_i \leftarrow w_i - \eta \operatorname{sign}(w_i)$$

✓ produce a large number of weights that are very near zero


- but almost never output weights that are exactly zero
- \rightarrow need to **threshold** away the very small weights to obtain a sparse solution

Gradient Descent

$$\min_x \underbrace{f(x)}_{\text{smooth}}$$

LOOP

- 1 find descent direction
- 2 descent



$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

easy to implement

$$\begin{aligned} x_{t+1} &= \arg \min_x \nabla f(x_t)^T (x - x_t) + \frac{1}{2\eta} \|x - x_t\|^2 \\ &= \arg \min_x \underbrace{f(x_t) + \nabla f(x_t)^T (x - x_t)}_{\text{linear approximation}} + \underbrace{\frac{1}{2\eta} \|x - x_t\|^2}_{\text{proximal term}} \end{aligned}$$

Proximal Gradient Descent

$$\min_{\mathbf{x}} \underbrace{f(\mathbf{x})}_{\text{smooth}} + \underbrace{g(\mathbf{x})}_{\text{nonsmooth}}$$

$$\begin{aligned} w_{t+1} &= \arg \min_w f(w_t) + \nabla f(w_t)'(w - w_t) + \frac{1}{2\eta} \|w - w_t\|^2 + g(w) \\ &= \underbrace{\arg \min_w g(w) + \frac{1}{2\eta} \|w - z_t\|^2}_{\text{proximal step}} \quad (\text{where } z_t = w_t - \eta \nabla f(w_t)) \end{aligned}$$

$$\text{prox}_{\eta g}(z) \equiv \arg \min_w \frac{1}{2} \|w - z\|^2 + \eta g(w)$$

$\text{prox}_{\eta g}(z)$ often has simple closed-form solution!


Proximal Step

Example (ℓ_2 regularizer)

$$\begin{aligned}\text{prox}_{\eta \frac{\lambda}{2} \|\cdot\|^2}(z) &= \arg \min_w \eta \frac{\lambda}{2} \|w\|^2 + \frac{1}{2} \|w - z\|^2 \\ &= \frac{1}{1 + \eta \lambda} \cdot z \quad (\text{shrinkage})\end{aligned}$$

Example (ℓ_1 regularizer)

$$\begin{aligned}\text{prox}_{\eta \lambda \|\cdot\|_1}(z) &= \arg \min_w \eta \lambda \|w\|_1 + \frac{1}{2} \|w - z\|^2 \\ &= [\text{sign}(z_i) \cdot \max(|z_i| - \eta \lambda, 0)] \quad (\text{soft-thresholding})\end{aligned}$$



Empirically, in Deep Learning: ℓ_1 vs ℓ_2

ℓ_1 regularization

- results in more zero parameters than ℓ_2 regularization
- better accuracy after pruning, but before retraining

ℓ_2 regularization

- higher accuracy than ℓ_1 regularization after retraining

Regularizing Network Weights: ℓ_0

ℓ_0 regularizer

$$R(\mathbf{w}) = \|\mathbf{w}\|_0$$

heuristic update procedure

- every n updates: set to zero all but the t elements of the parameter vector with the largest magnitude

Deep Compression

combine pruning with other techniques for further compression


(Han et al. 2016)

Deep compression: Compressing deep neural network with pruning, trained quantization and Huffman coding

- 1 pruning
 - train → prune small-weight connections → retrain
- 2 **weight sharing**
 - cluster the weights (k -means clustering)
 - generate **codebook**
 - multiple weights share the same codebook
 - retrain codebook
- 3 **Huffman coding** of clustered weights

Experimental Results

Accuracy vs compression rate




Problem with Pruning

- connections may be wrongly pruned
- once pruned, connections gone forever

(Guo et al., 2017)

Dynamic Network Surgery for Efficient DNNs

splicing: enable recovery of important connections



Which Connections to Prune or Splice?

cannot just look at the weight magnitude

- associate an **importance measure** t_i to each weight w_i
 - $t_i = 0$: unimportant \rightarrow prune
 - $t_i = 1$: important \rightarrow keep
- use thresholds a, b and update t_i as

$$t_i \leftarrow \begin{cases} 0 & a > |w_i| & (\text{prune}) \\ t_i & a \leq |w_i| < b \\ 1 & b \leq |w_i| & (\text{splice}) \end{cases}$$

Dynamic Network Surgery

procedure

- ➊ forward propagation with $\mathbf{w} \odot \mathbf{t}$
- ➋ backward propagation
- ➌ update t_i
 - but only with probability $\sigma(\text{iter})$
 - slow down pruning and splicing frequencies
- ➍ update w_i (including those with zero t_i s)

$$w_i = w_i - \eta \frac{\partial}{\partial (w_i t_i)} \ell(\mathbf{w} \odot \mathbf{t})$$

Experimental Results

AlexNet

Layer	Params.	Params.% [9]	Params.% (Ours)
conv1	35K	~ 84%	53.8%
conv2	307K	~ 38%	40.6%
conv3	885K	~ 35%	29.0%
conv4	664K	~ 37%	32.3%
conv5	443K	~ 37%	32.5%
fc1	38M	~ 9%	3.7%
fc2	17M	~ 9%	6.6%
fc3	4M	~ 25%	4.6%
Total	61M	~ 11%	5.7%

[from (Guo et al., 2017)]

- reduces model complexity, while the prediction error rate does not increase


Another Problem

- pruning usually done layer by layer
- weights may appear unimportant in an early layer
- but may contribute significantly to important units in later layers

Neuron Importance Score Propagation (NISP)

(Yu et al., 2018)

NISP: Pruning Networks using Neuron Importance Score Propagation



- 1 obtain **importance scores** of units in the last hidden layer (using some feature ranking algorithm)
- 2 **propagate importance scores**
- 3 prune less important units
- 4 fine-tuning

Pros and Cons of Weight Pruning

pros


- fine-grained, most flexible
- can lead to high compression rate

cons

- standard ML frameworks do not support efficient operation on sparse weight tensors
- require special libraries or hardware for fast inference

Limitations of Weight Pruning in CNN

AlexNet



[from (Krizhevsky, 2012)]


- weights removed mostly come from the fully-connected layers
 - computations in these layers are cheap
 - resultant time reduction is insignificant
 - many new deep networks use fewer fully-connected layers
- convolution is more costly → runtime during prediction is dominated by evaluation of convolution layers

Limitations of Weight Pruning in RNN

- In RNN, the basic structures interweave with each other

Example (LSTM)

- one hidden unit is associated with entries in all eight matrices, and entries in the corresponding entries in the next layer



- independently removing these structures can result in mismatch of their dimensions and then inducing invalid recurrent units

Structured Pruning

Structured Pruning in CNN

prune filters/channels directly

pros

- does not introduce sparsity → does not require sparse libraries or any specialized hardware
- easy to set a target number of filters to be pruned

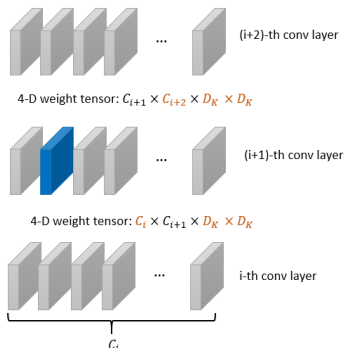
cons

- coarse-grained, less flexible

Pruning Filters: Based on Weights

(Li et al., 2017)

Pruning filters for efficient ConvNets



- 1 measure importance of each filter
 - e.g., sum of absolute kernel weights
- 2 prune some filters with **smallest** importance
 - remove corresponding feature map
 - kernels that apply on the removed feature maps from the filters of the next convolutional layer are also removed
- 3 after pruning, retrain the network

Pruning Filters: Based on Other Criteria

(Molchanov et al., 2017)

Pruning convolutional neural networks for resource efficient transfer learning

find a network (with parameter \mathbf{w}') such that

- it is **small**
- its **performance** (e.g., negative log-likelihood) is close to that of the original network (with parameter \mathbf{w})

$$\min_{\mathbf{w}'} |\ell(\mathbf{w}') - \ell(\mathbf{w})| : \|\mathbf{w}'\|_0 \leq B$$

Procedure

$$\min_{\mathbf{w}'} |\ell(\mathbf{w}') - \ell(\mathbf{w})| : \|\mathbf{w}'\|_0 \leq B$$

- ① evaluate **importance** of feature maps
 - weight magnitude (assume that small weights are less important)
 - hidden unit activation
 - estimated change in loss after removing weights (use Taylor expansion)
 - based on **batch normalization**
- ② remove the least important feature map
- ③ fine-tuning
- ④ repeat until $\|\mathbf{w}'\|_0 \leq B$

Batch Normalization (BN)

- normalizes activations for mini-batch $B = \{x_1, \dots, x_m\}$ (usually done before nonlinearity)
- \rightarrow helps training of deep networks

batch normalization

mini-batch mean: $\mu_B \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$

mini-batch variance: $\sigma_B^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$

normalize: $\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$

scale and shift: $y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i)$

- γ (scaling) and β are learnable parameters

Batch Normalization in CNN

- input X : size $B \times X \times Y \times C$

- 1 $\mu_c = \frac{1}{BXY} \sum_{bxy} x_{bxyz}$

- 2 $\sigma_c^2 = \frac{1}{BXY} \sum_{bxy} (x_{bxyz} - \mu_c)^2$

- 3 output

$$y_{bxyz} = \left(\frac{x_{bxyz} - \mu_c}{\sqrt{\sigma_c^2 + \epsilon}} \right) \cdot \gamma_c + \beta_c$$

Pruning Channels: Network Slimming

$$y_{bxyz} = \left(\frac{x_{bxyz} - \mu_c}{\sqrt{\sigma_c^2 + \epsilon}} \right) \cdot \gamma_c + \beta_c$$


(Liu et al., 2017)

Learning Efficient Convolutional Networks through **Network Slimming**

idea

- magnitude of $\gamma_c \rightarrow$ importance of channel c
- to select channels
 - \rightarrow enforcing **sparsity** on these scaling parameters
 - \rightarrow imposes ℓ_1 regularization

Network Slimming Procedure



- ① jointly train the network weights and BN parameters, with ℓ_1 regularization imposed on γ_c 's
 - ✓ easy to implement without introducing any change to existing CNN architectures
- ② prune channels with small factors
- ③ fine-tuning
- ④ repeat

Structured Sparsity

Features often have intrinsic **structures**

Example (group sparsity)

a categorical feature (e.g., nationality) is represented by a **group** of dummy binary features (european/american/chinese)



- select/drop whole groups
- **group sparsity**: a few groups are selected

Group Lasso Regularizer

- an extension of the lasso for feature selection on (predefined) groups of features
- given a number of groups G_i 's
- $r(w) = \sum_i \|w_{G_i}\|_2$
 - take the ℓ_2 norm of all the groups
 - ℓ_1 norm of the above (sparsity on groups)
- reduces to the lasso if groups have cardinality one

Using Group Lasso in CNNs

(Wen et al, 2016)

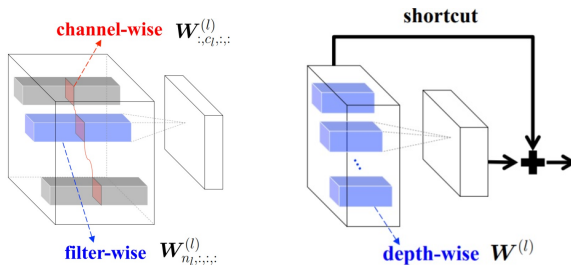
Learning structured sparsity in deep neural networks

minimize loss + sparsity-inducing regularizer

$$\min E_D(W) + \lambda R(W) + \lambda_g \sum_{l=1}^L R_g(W^{(l)})$$

- $E_D(W)$: loss on data
- $R(W)$: ℓ_2 regularizer
- R_g : **group lasso** regularizer

Remove Less Important Structures in CNN



- define each **channel** as a group → remove less important channels
- define each **filter** as a group → remove less important filters
- define each **depth** as a group → reduce the depth

Group Lasso in LSTM

(Wen et al., 2018)

Learning intrinsic sparse structures within long short-term memory

$$\mathbf{i}_t = \sigma(\mathbf{x}_t \mathbf{W}_{xi} + \mathbf{h}_{t-1} \mathbf{W}_{hi} + \mathbf{b}_i)$$


$$\mathbf{f}_t = \sigma(\mathbf{x}_t \mathbf{W}_{xf} + \mathbf{h}_{t-1} \mathbf{W}_{hf} + \mathbf{b}_f)$$

$$\mathbf{a}_t = \tanh(\mathbf{x}_t \mathbf{W}_{xa} + \mathbf{h}_{t-1} \mathbf{W}_{ha} + \mathbf{b}_a)$$

$$\mathbf{o}_t = \sigma(\mathbf{x}_t \mathbf{W}_{xo} + \mathbf{h}_{t-1} \mathbf{W}_{ho} + \mathbf{b}_o)$$

$$\mathbf{c}_t = \mathbf{i}_t \odot \mathbf{a}_t + \mathbf{f}_t \odot \mathbf{c}_{t-1}$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$




- prune a hidden state \rightarrow prune corresponding row
- prune input-gate/forget-gate/cell/output-gate \rightarrow prune corresponding column

Other Structured Sparsity Regularizers

Example

with unknown groups, can we obtain the group structure automatically?



feature grouping

- group relevant and highly correlated features

OSCAR

(Bondell & Reich, 2008)

Simultaneous regression shrinkage, variable selection, and supervised clustering of predictors with OSCAR

Octagonal **S**hrinkage and **C**lustering **A**lgorithm for **R**egression

$$\min_w \|y - Xw\|^2 + \lambda_1 \underbrace{\|w\|_1}_{\text{sparsity}} + \lambda_2 \underbrace{\sum_{i < j} \max\{|w_i|, |w_j|\}}_{\text{grouping}}$$

- encourages **equality** of coefficients for correlated features → automatic grouping

(Zhang et al., 2018)


Learning to share: Simultaneous parameter tying and sparsification in deep learning

- GrOWL (group ordered weighted ℓ_1)

Topology of Pruned Networks

Pruning


- 1 train a network (with full-precision weights)
- 2 prune unimportant weights



- 3 retrain the pruned network
- 4 (repeat...)

(deep networks) pruning/retraining \leftrightarrow (biological networks)
weakening/strengthening



Networks and Power Law



- distribution of node degree follows the **power law**



$$p(x) \propto x^{-\alpha}$$

- e.g., internet, collaboration, web, power grid, ...



Scale-Free Brain Functional Networks

- distribution of functional connections also follows the **power law** [Eguiluz et al, 2005]



Are artificial neural networks scale-free?

(Lu and Kwok, 2018)

Power Law in Sparsified Deep Neural Networks


MLP on MNIST

MLP: input - 1024FC - 1024FC - 10softmax


- truncated power-law distribution (TPL) [Kolyukhin and Torabi, 2013]

$$p(x) = \frac{1 - \alpha}{x_{\max}^{1-\alpha} - x_{\min}^{1-\alpha}} x^{-\alpha}$$


- $[x_{\min}, x_{\max}]$: range over which the power law is valid



(a) input.




(b) fc1.




(c) fc2.

CNN on MNIST and CIFAR-10


input - 32C5 - MP2 - 32C5 - MP2 - 256FC - 10softmax



(a) conv1.








(b) conv2.




(c) fc3.

input - 32C3 - 32C3 - MP2 - 64C3 - 64C3 - MP2 - 512FC - 10softmax




Larger CNNs on ImageNet

AlexNet



VGG-16



Preferential Attachment


- power law distributions can originate from the process of **preferential attachment** [Barabasi and Albert, 1999]

preferential attachment

- network evolves over time
 - new nodes are more likely to link to nodes that already have high degrees
 - **rich gets richer**
- when a **new edge** is added to the network, existing nodes with high degrees are more likely to connect to each other

Do we have preferential attachment in artificial neural networks?

Continual Learning



Time $t = 0$

- task A: MNIST image classification
- a dense network is trained, pruned and re-trained

Time $t = 1$

- new task B: MNIST image classification, but pixels inside a central square of each image are permuted
- add back pruned connections, re-train on task B
- prune, and re-train on task B

Number of Connections Created


consider two nodes in consecutive layers

- one with degree d_1 , the other with degree d_2


if there is preferential attachment ...

number of connections created between these two nodes

$$\propto d_1 d_2$$



(a) input-to-fc1.




(b) fc1-to-fc2.

Increase in Node Degree


- $d_i(t)$: degree of node i at time t

increase in node degree


$$\frac{dd_i(t)}{dt} \propto d_i(0)$$



(a) input.



(b) fc1.




(c) fc2.


Degree Distribution

degree distribution before/after adding new task


If initial degree distribution follows the power law, the final degree distribution follows the same power law



(a) input.




(b) fc1.




(c) fc2.

Progressive Network



- ① train on task A using the left part of the network, prune, re-train
- ② fix the parameters
- ③ train on task B: add the right part of the network
 - standard progressive network: $h_2^{(2)}$ is fully connected to $h_1^{(1)}$, ...
 - using **preferential attachment**: probability of connecting to a node in $h_1^{(1)}$ is proportional to its degree

Results on MNIST



- preferential attachment always outperforms random attachment
- 5% lateral connections to task A
 - preferential attachment has significantly better accuracy than using separate networks
- 20% lateral connections
 - preferential attachment has similar or even better accuracies than the dense progressive neural network