

**facebook**

Artificial Intelligence Research

# Using Neural Networks for Modeling and Representing Natural Languages : *Introduction to Neural Networks*

Tomas Mikolov, Piotr Bojanowski, Facebook AI Research

Monday, July 22nd 2019

# Structure of the tutorial

- Motivation
- Basic machine learning applied to natural language
- Introduction to neural networks
- Distributed representations of words
- Efficient text classification
- Neural network based language models
- Future research
- Resources

# Introduction

- Text processing is the core business of internet companies today (Facebook, Google, Twitter, Baidu, Yahoo, ...)
- Machine learning and natural language processing techniques are applied to big datasets to improve many tasks:
  - Search
  - Ranking
  - Spam detection, fake news detection, ads recommendation, email categorization, machine translation, speech recognition, ...

# Introduction

- This tutorial introduces artificial neural networks applied to text problems
- The focus is on understanding the core ideas: how do artificial neural networks work, what they can and cannot do, what is deep learning
- Overview of some interesting results that have been already achieved

# Basic Machine Learning for NLP

# Basic machine learning applied to NLP

- Before we start talking about neural networks, basic techniques will be briefly mentioned
- Neural networks are closely related to other basic machine learning techniques
- To avoid re-discovery of the wheel, it is important to know the basic concepts first
- Finally: while the basic techniques are often trivial, it is very hard to improve upon them (and many fancy techniques fail to do so!)

# Basic machine learning applied to NLP

- N-grams
  - Word classes
  - Bag-of-words representations
- 
- Logistic regression
  - Support vector machines

# N-grams

- Standard approach to language modeling

- Task: compute probability of a sentence  $W$

$$P(W) = \prod_i P(w_i | w_1 \dots w_{i-1})$$

- Often simplified to trigrams:

$$P(W) = \prod_i P(w_i | w_{i-2}, w_{i-1})$$



# N-grams: example

$$P(\text{"this is a sentence"}) = P(\text{this}) \times P(\text{is}|\text{this}) \times P(\text{a}|\text{this, is}) \times P(\text{sentence}|\text{is, a})$$

- The probabilities are estimated from counts in some (large) text corpus:

$$P(a|\text{this, is}) = \frac{C(\text{this is a})}{C(\text{this is})}$$

- Smoothing is used to redistribute probability to unseen events (this avoids zero probabilities)

*A Bit of Progress in Language Modeling* (Goodman, 2001)

# Word classes

- One of the most successful NLP concepts in practice
- Similar words should share parameters, which leads to generalization
- Example:

$Class_1 = (yellow, green, blue, red)$   
 $Class_2 = (Italy, Germany, France, Spain)$

- Usually, each vocabulary word is mapped to a single class (similar words share the same class)

# Word classes

- There are many ways how to compute the classes – usually, it is assumed that similar words appear in similar contexts
- Instead of using just counts of words, we can use counts of classes, which leads to generalization (better performance on novel data)

*Class-based  $n$ -gram models of natural language* (Brown, 1992)

# One-hot representations

- Simple way how to encode discrete concepts, such as words

Example:

```
vocabulary = (Monday, Tuesday, is, a, today)
```

```
Monday    = [1 0 0 0 0]
```

```
Tuesday   = [0 1 0 0 0]
```

```
is        = [0 0 1 0 0]
```

```
a         = [0 0 0 1 0]
```

```
today     = [0 0 0 0 1]
```

Also known as 1-of-N coding (N would be the size of the vocabulary)

# Bag-of-words representations

- Sum of one-hot codes
- Ignores order of words

Example:

```
vocabulary = (Monday, Tuesday, is, a, today)
```

```
Monday Monday = [2 0 0 0 0]
```

```
today is a Monday = [1 0 1 1 1]
```


```
today is a Tuesday = [0 1 1 1 1]
```

```
is a Monday today = [1 0 1 1 1]
```

Can be extended to bag-of-N-grams to capture local ordering of words

# Logistic regression

- Simple machine learning technique to perform classification
- Input is a vector of features, output is usually one (binary classification) or many (multinomial distribution)



- The weight matrix (or vector) directly connects inputs and output(s)

# Logistic regression

- Can be trained by stochastic gradient descent, and can be seen as a neural network without any hidden layers (will be described later)
- Also called maximum entropy model in the NLP community
- Example C code for toy problems available at:  
<http://ai.stanford.edu/~ajoulin/code/nn.zip>  
(joint work with Armand Joulin; includes code for logistic regression, feedforward and recurrent neural networks)

# Support vector machines

- Another popular way how to perform classification, very similar to logistic regression
- Tries to maximize margin between the classes:




Figure from Wikipedia

- Used to be popular in part because of existence of open-source packages: *libsvm*, *svmtorch*, *svmlight*



# Basic machine learning: summary

Main statistical tools for NLP:

- Count-based models: N-grams, bag-of-words
- Word classes
- Unsupervised dimensionality reduction: PCA
- Unsupervised clustering: K-means
- Supervised classification: logistic regression, SVMs

# Neural Networks


# Introduction to neural networks

- Motivation
- Architecture of neural networks: neurons, layers, synapses
- Activation function
- Objective function
- Training: stochastic gradient descent, backpropagation, learning rate, regularization
- Intuitive explanation of deep learning


# Neural networks: motivation

- The main motivation in NLP is to come up with more precise way how to represent and model words, documents and language
- There is nothing that neural networks can do in NLP that the basic techniques completely fail at
- But: the victory in competitions goes to the best, thus few percent gain in accuracy is important!


# Neuron (perceptron)




# Neuron (perceptron)



# Neuron (perceptron)




# Neuron (perceptron)






# Neuron (perceptron)



# Neuron (perceptron)




# Neuron (perceptron)

- It should be noted that the perceptron model is quite different from the biological neurons (those communicate by sending spike signals at various frequencies)
- The learning in biological neurons seems to be also quite different
- It would be better to think of artificial neural networks as non-linear projections of data

# Activation function


- In the previous example, we used  $\max(0, \text{value})$ : this is nowadays referred to as “rectified linear unit” (ReLU)
- Many other non-linear functions can be used
- Other common ones: sigmoid, tanh



Sigmoid function,  
Figure from Wikipedia

# Activation function

- The most important property is the non-linearity
- Example: XOR problem
  - There is no linear classifier that can solve this problem:



# Non-linearity: example

Intuitive NLP example:

- Input: bag-of-words
- Output: binary classification (for example, positive / negative sentiment)
- Input: “*the idea was not bad*”
- Non-linear classifier can learn that “*not*” and “*bad*” next to each other mean something else than “*not*” or “*bad*” itself
- “*not bad*”  $\neq$  “*not*” + “*bad*”

# Activation function

- The non-linearity is a crucial concept that gives neural networks more representational power compared to some other techniques (linear SVMs, logistic regression)
- Without the non-linearity, it is not possible to model certain combinations of features (like Boolean XOR function), unless we do manual feature engineering


# Hidden layer

- Hidden layer represents learned non-linear combination of input features (this is different than SVMs with non-linear kernels that are not learned)
- Inputs to a hidden layer can be outputs of a previous hidden layer



# Hidden layer


- With a hidden layer, we can solve the XOR problem:
  1. some neurons in the hidden layer will activate only for some combination of the input features
  2. the output layer can represent a combination of the activations of the hidden neurons



# Hidden layer

- Neural net with one hidden layer is a universal approximator: it can represent any function up to arbitrary precision
- However, not all functions can be represented *efficiently* with a single hidden layer – we shall see that in the deep learning section

# Neural network layers



INPUT LAYER

HIDDEN LAYER

OUTPUT LAYER

# Objective function

- Objective function defines how well does the neural network perform some task
- The goal of training is to adapt the weights so that the objective function is optimized (maximized / minimized)
- Example: classification accuracy, reconstruction error

# Unsupervised / supervised training

- When the objective is to model the input data, the training is called unsupervised
- An example is auto-encoder: the objective function is to reconstruct the input features at the output layer (by performing some kind of compression when going through the hidden layers)
- Supervised training means that we have additional labels for the input vectors, and the objective is usually to perform classification


# Training of neural networks

- There are many ways how to train neural networks
- The most widely used and successful in practice is stochastic gradient descent (SGD)
- Many algorithms are presented as superior to SGD, but when properly compared, the improvements are not easy to achieve

# Training of neural networks

Forward pass:

- Input signal is presented first
- Hidden layer state is computed (vector times matrix operation and non-linear activation)
- Outputs are computed (vectors times matrix operation and usually non-linear activation)



$W$ : input weights

Activation function:

$\max(0, \text{value})$

$I$ : input signal

$\text{Output} = \max(0, I \cdot W)$

# Training of neural networks - SGD

Intuitive explanation of stochastic gradient descent:

- The input feature vector is used to compute the output vector during the forward pass
- The target vector represents the desired output vector (in case of classification it uses one-hot coding)
- We change the weights a little bit so that next time the same input vector is presented, the output vector will be closer to the target vector



# Training of neural networks - SGD

Intuitive explanation of stochastic gradient descent:

- The input feature vector is used to compute the output vector during the forward pass
- The target vector represents the desired output vector (in case of classification it uses one-hot coding)
- We change the weights a little bit so that next time the same input vector is presented, the output vector will be closer to the target vector



# Training of neural networks - SGD

Intuitive explanation of stochastic gradient descent:

- The input feature vector is used to compute the output vector during the forward pass
- The target vector represents the desired output vector (in case of classification it uses one-hot coding)
- We change the weights a little bit so that next time the same input vector is presented, the output vector will be closer to the target vector

# Backpropagation

- To train the network, we need to compute gradient of the error
- The gradients are sent back using the same weights that were used in the forward pass
- 
- Simplified graphical representation:



# Training of neural networks – learning rate

- Learning rate controls how quickly we change the weights: too little value will result in long training time, too high value will erase what was learned previously
- In practice, we start with a high learning rate and reduce it during the training
- The starting learning rate and how quickly it gets reduced can affect the resulting performance in a great way: you have to tune this!

# Training of neural networks – training epochs

- Several training epochs over the training data are often performed
- Usually, the training is finished when performance on held-out (validation) data does not improve
- Held-out data is used only for verification of the performance, the network is not trained on these examples

# Regularization

- As the network is trained, it can overfit the training data
  - Overfitting: very good performance on training and bad performance on test data
- The network can “memorize” the training data: often, it will contain high weights that are used to model only some small subset of the data
- We can force the weights to stay small (close to zero) during training to reduce this problem (L1 & L2 regularization)

# Training of neural networks: summary

- Stochastic gradient descent and backpropagation are usually good choices for training
- The representational power of neural networks comes from non-linear hidden layer(s)

# What training typically does not do

- Choice of the hyper-parameters has to be done manually:
  - Type of activation function
  - Choice of architecture (how many hidden layers, their sizes)
  - Learning rate, number of training epochs
  - What features are presented at the input layer
  - How to regularize
- 
- It may seem complicated at first - the best way to start is to re-use some existing setup and try your own modifications.



# Neural networks and logistic regression

- Neural networks can do everything logistic regression can do (proof: the hidden layer can simply copy inputs)
- Logistic regression is in many cases computationally much more efficient
- Thus, we could use both jointly to be efficient: will be shown later

# Deep learning


- Deep model architecture is about having more than one computational step (hidden layer) in the model
- Deep learning aims to learn patterns that cannot be learned efficiently with shallow models (one or zero hidden layers)
  - Should result in better generalization (performance on the test set)

# Deep learning

- But: it was previously stated that one hidden layer is enough to represent any function
- Why would we need more hidden layers then?

# Deep learning

- The crucial part to understand deep learning is the *efficiency*
- The “*universal approximator*” argument says nothing else than that a neural net with non-linearities can work as a look-up table to represent any function: some neurons can activate only for some specific range of input values



# Deep learning

- Look-up table is not efficient: for certain functions, we would need exponentially many hidden units with increasing size of the input layer
- Example of function that is difficult to represent: parity function (N bits at input, output is 1 if the number of active input bits is odd)


*Perceptrons*, Minsky & Papert 1969

# Deep learning

- Having hidden layers exponentially larger than is necessary is bad: too many parameters to learn
- If we cannot compactly represent patterns, we have to memorize them: we need (possibly exponentially) more training examples than is necessary

# Deep learning

- Whenever we try to learn a complex function that is a composition of simpler functions, it may be beneficial to use a deep architecture



# Deep learning

- Historically, deep learning was assumed to be impossible to achieve by using SGD + backpropagation
- Since ~2010 there was a lot of progress in tasks that contain signals that are very compositional (speech, vision) and where large amount of training data is available



# Deep learning

- Deep learning is still an open research problem
- Many deep models have been proposed that do not learn anything else than a shallow (one hidden layer) model can learn: beware the hype!
- Not everything labeled “deep” is a successful example of deep learning
  - Having a deep architecture is not sufficient to learn complex patterns

# Neural networks and deep learning: summary

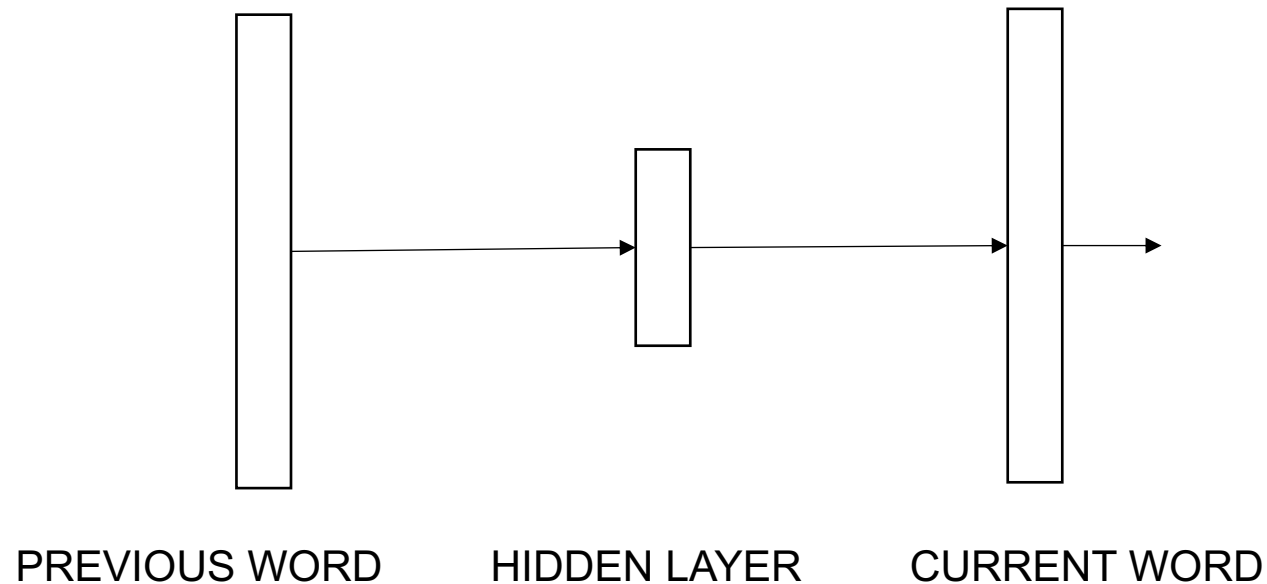
- Neural networks are a basic machine learning technique, can be seen as non-linear projections of the input feature vectors
- Start with SGD and backpropagation for training
- Deep learning can be useful for learning complex patterns in data, especially in vision & speech

# Distributed Word Representations

# Distributed representations of words

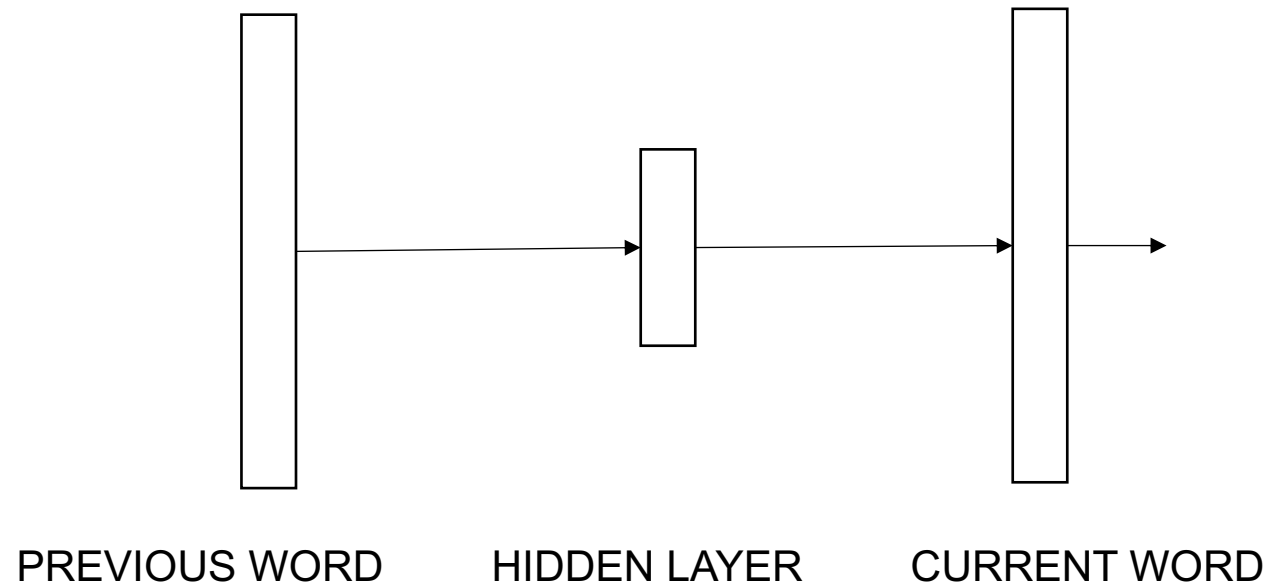
- Vector representation of words computed using neural networks
- Linguistic regularities in the word vector space
- Evaluation of performance
- Application to machine translation
- Recent progress

# A very basic neural network applied to NLP



- Bigram neural language model
- Previous word is used to predict the current word by going through hidden layer (classifier with as many outputs as there are words in the vocabulary)

# A very basic neural network applied to NLP



- The input is encoded as one-hot
- The model will learn compressed, continuous representation of words (usually the matrix of weights between the input and hidden layer)

# Word vectors

- We call the vectors in the matrix between the input and hidden layer *word vectors* (also known as *word embeddings*)
- Each word is represented by a real valued vector in N-dimensional space (usually  $N = 50 - 1000$ )
- The word vectors have some similar properties to word classes; however, many degrees of similarity are captured
  - *Paris* is similar to *Berlin*, but also to *France*

# Word vectors

- Word vectors can be used as features in many NLP tasks
  - *Natural Language Processing (Almost) from Scratch*, Collobert et al 2011
- Pre-trained word vectors provide generalization for systems trained with limited amount of supervised data
- Complex model architectures can be used to learn the word vectors
  - Neural language model with multi-task learning as in (Collobert & Weston 2008)



# Word vectors

- Many architectures were proposed for training the word vectors, some labeled “deep”
- Do we really need deep learning here?
- Do we know how to apply deep learning to this task?

# Word vectors

- We need some measurable way how to compare word vectors trained using different architectures
- The comparison is tricky: people mostly publish just their pre-trained word vectors, use different datasets both for training and evaluation...
- Conclusions based on experiments with different datasets are difficult to make

# Word vectors - evaluation

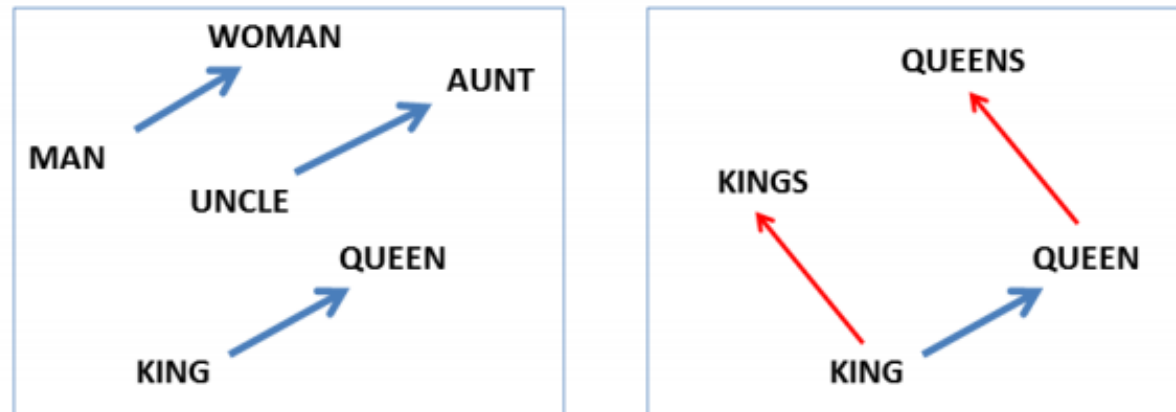
Popular datasets like WS353 (word similarity, 353 word pairs with human judgements of similarity) have several drawbacks:

- Tiny size, no heldout / test data split
- Performance is heavily biased by choice of the training data, less so by the architecture of the model itself

*Placing Search in Context: The Concept Revisited* (Finkelstein et al, 2002)

# Word vectors – linguistic regularities

- It was shown that word vectors capture many linguistic properties (gender, tense, plurality, even semantic concepts like “capital city of”)
- We can do nearest neighbor search around result of vector operation “King – man + woman” and obtain “Queen”



*Linguistic regularities (Mikolov et al, 2013)*

Mikolov et al,

# Word vectors – datasets for evaluation

Microsoft Research dataset with 8K “analogies” - examples:

- good: better    rough: \_\_\_\_
- good: best     rough: \_\_\_\_
- better: best    rougher: \_\_\_\_
- year: years    law: \_\_\_\_
- see: saw        return: \_\_\_\_

*Linguistic regularities in continuous space word representations* (Mikolov et al, 2013)

# Word vectors – datasets for evaluation

Google dataset, almost 20K questions:

- Athens:Greece                      Oslo: \_\_\_\_\_
- Angola:kwanza                      Iran: \_\_\_\_\_
- brother:sister                      grandson: \_\_\_\_\_
- possibly:impossibly                      ethical: \_\_\_\_\_
- walking:walked                      swimming: \_\_\_\_\_

*Efficient estimation of word representations in vector space* (Mikolov et al, 2013)

# Word vectors – datasets for evaluation

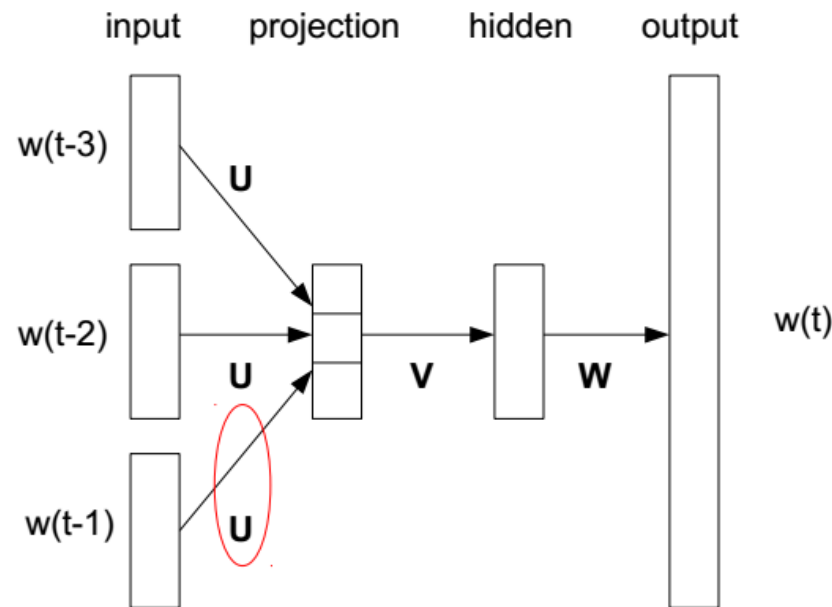
Google phrase-based dataset, focuses on semantics:

- New York:New York Times      Baltimore: \_\_\_\_
- Boston:Boston Bruins      Montreal: \_\_\_\_
- Detroit:Detroit Pistons      Toronto: \_\_\_\_
- Austria:Austrian Airlines      Spain: \_\_\_\_
- Steve Ballmer:Microsoft      Larry Page: \_\_\_\_

*Distributed Representations of Words and Phrases and their Compositionality*  
(Mikolov et al, 2013)


# Word vectors – various architectures

- Neural net based word vectors were traditionally trained as part of full neural network language model (Bengio et al, 2003)
- This models consists of an input layer, projection layer, hidden layer and output layer (will be discussed in detail in the next section)






# Word vectors – various architectures



- We can extend the bigram NNLM by adding more context

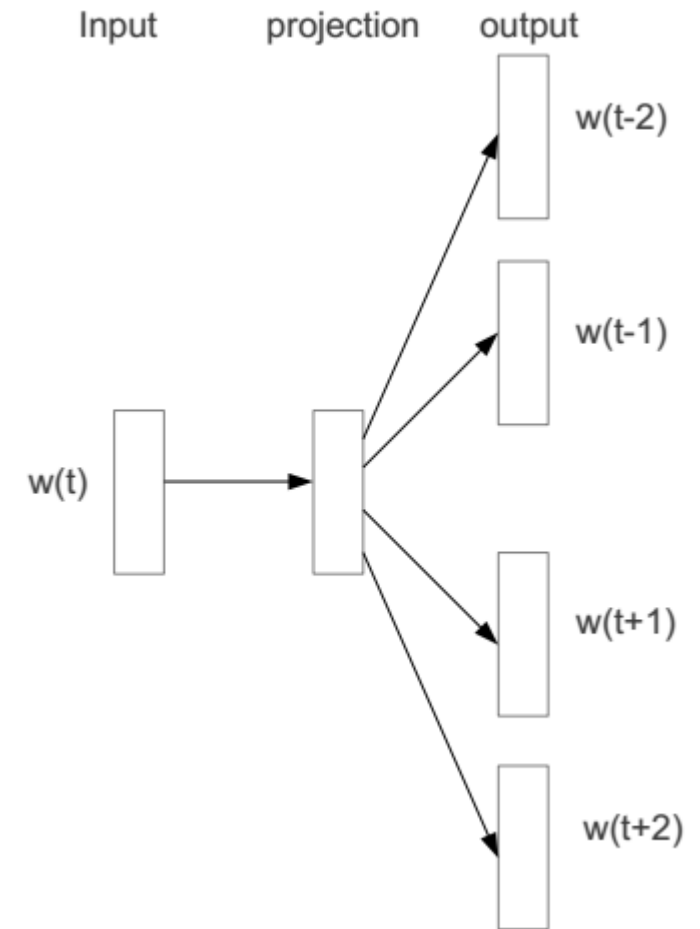
# Word vectors – various architectures

- The ‘continuous bag-of-words model’ (CBOW) adds inputs from words within short window to predict the current word
- The weights for different positions are shared
- Computationally much more efficient than normal NNLM (but cannot model n-grams)
- The hidden layer is linear (no activation)



# Word vectors – various architectures

- We can reformulate the CBOW model by predicting surrounding words using the current word
- This architecture is called skip-gram
- Similar performance to CBOW after convergence



# Word vectors - training

- SGD + backpropagation
- Approximation of very large softmax in the output layer – can easily be in order of millions of outputs (computationally expensive):
  1. Hierarchical softmax (will be described in the LM section)
  2. Negative sampling

# Word vectors – negative sampling

- Instead of propagating signal from the hidden layer to the whole output layer, only the output that represents the positive class + few randomly sampled outputs are evaluated
- These outputs are treated as independent logistic regression classifiers
- This makes the training speed independent on the vocabulary size

# Word vectors - subsampling

- It is useful to sub-sample the frequent words ('the', 'is', 'a', ...) during training (discard with probability proportional to the word frequency)
- Non-linearity does not seem to improve performance of these models, thus the hidden layer does not use activation function

# Word vectors – comparison of performance

<i>Model</i>	<i>Vector Dimensionality</i>	<i>Training Words</i>	<i>Training Time</i>	<i>Accuracy [%]</i>
Collobert NNLM	50	660M	2 months	11
Turian NNLM	200	37M	few weeks	2
Mnih NNLM	100	37M	7 days	9
Mikolov RNNLM	640	320M	weeks	25
Huang NNLM	50	990M	weeks	13
Skip-gram (hier.s.)	1000	6B	hours	66
CBOW (negative)	300	1.5B	<b>minutes</b>	<b>72</b>

- Google 20K questions dataset (word based, both syntax and semantics)
- Going from weeks of training to minutes while improving accuracy!

# Word vectors – scaling up

- The choice of training corpus is usually more important than the choice of the technique itself
- Low computational complexity is crucial
- Optimized code published as word2vec project:  
<https://code.google.com/p/word2vec/>  
<https://github.com/tmikolov/word2vec>



# Word vectors – nearest neighbors


	<b>Redmond</b>	<b>Havel</b>	<b>graffiti</b>	<b>capitulate</b>
<b>Collobert NNLM</b>	conyers lubbock keene	plauen dzerzhinsky osterreich	cheesecake gossip dioramas	abdicate accede rearm
<b>Turian NNLM</b>	McCarthy Alston Cousins	Jewell Arzu Ovitz	gunfire emotion impunity	- - -
<b>Mnih NNLM</b>	Podhurst Harlang Agarwal	Pontiff Pinochet Rodionov	anaesthetics monkeys Jews	Mavericks planning hesitated
<b>Skip-gram (phrases)</b>	Redmond Wash. Redmond Washington Microsoft	Vaclav Havel president Vaclav Havel Velvet Revolution	spray paint grafitti taggers	capitulation capitulated capitulating

- More training data helps the quality a lot!

# Word vectors – more examples

<i>Expression</i>	<i>Nearest token</i>
Paris - France + Italy	Rome
bigger - big + cold	colder
sushi - Japan + Germany	bratwurst
Cu - copper + gold	Au
Windows - Microsoft + Google	Android
Montreal Canadiens - Montreal + Toronto	Toronto Maple Leafs

# Word vectors – visualization using PCA



# Vectors: from words to phrases

- Linguistically, *New York* or *Air Canada* should be treated as units
- Phrases can be constructed using mutual information criterion between words
- Pre-process the training data and rewrite all phrases as single tokens, such as *New\_York* and *Air\_Canada*

# Sentence-level representations

- To obtain sentence level representations, we can add unique tokens to the data, one for each sentence (or short document)
- These tokens are trained in a similar way like other words in the skip-gram or CBOW models, just using unlimited context window (within the sentence boundaries)

Example:

SID\_\_1 We think this was not the best way ...

SID\_\_2 Another reason was to ...



# Sentence-level representations

- The sentence representations can be further used in classifiers (logistic regression, SVM, or neural network)
- Needs to be trained for many epochs; good results on sentiment analysis tasks

*Distributed Representations of Sentences and Documents (Le et al, 2014)*

# Translation of words using vector spaces

- Natural languages describe the same concepts – dogs have four legs, the sky is blue everywhere in the world




# Translation of words using vector spaces

- It could be enough to learn a mapping between the vector spaces to perform basic translation
- We tried to start with small existing dictionary (5K most frequent words), and tried to translate the remaining words



# English to Spanish: translation of words



Spanish word	Computed English Translations	Dictionary Entry
emociones	emotions emotion feelings	emotions
protegida	wetland undevelopable protected	protected
imperio	dictatorship imperialism tyranny	empire
destacaron	highlighted emphasized emphasised	highlighted

- The results are surprisingly accurate, especially with models trained on a lot of data

# Translation of words and phrases

- We could reach above 90% accuracy for the most confident translations
- This technique is useful when monolingual data is plentiful and bilingual data is rare (internet slang words, distant language pairs, ...)

*Exploiting similarities among languages for machine translation* (Mikolov et al, 2013)

# Distributed word representations: summary

- Simple model architectures seem to work the best
- Parameter tuning is still a bit of an art: context size, dimensionality, training algorithm, number of training epochs, ...
- Large text corpora are crucial for good performance (some links will be given in the Resources section)
- Current state of the art for word representations is the fastText project (see next session)